

Keras -- MLPs on MNIST

In [0]:

```
# if you keras is not using tensorflow as backend set "KERAS_BACKEND=tensorflow" use this command
from keras.utils import np_utils
from keras.datasets import mnist
import seaborn as sns
from keras.initializers import RandomNormal
```

In [0]:

```
%matplotlib inline
import matplotlib.pyplot as plt
import numpy as np
import time
# https://gist.github.com/greydanus/f6eee59eaf1d90fcb3b534a25362cea4
# https://stackoverflow.com/a/14434334
# this function is used to update the plots for each epoch and error
def plt_dynamic(x, vy, ty, ax, colors=['b']):
    ax.plot(x, vy, 'b', label="Validation Loss")
    ax.plot(x, ty, 'r', label="Train Loss")
    plt.legend()
    plt.grid()
    fig.canvas.draw()
```

In [0]:

```
# the data, shuffled and split between train and test sets
(X_train, y_train), (X_test, y_test) = mnist.load_data()
```

In [80]:

```
print("Number of training examples :", X_train.shape[0], "and each image is of shape (%d, %d)"%(X_train.shape[1], X_train.shape[2]))
print("Number of training examples :", X_test.shape[0], "and each image is of shape (%d, %d)"%(X_test.shape[1], X_test.shape[2]))
```

Number of training examples : 60000 and each image is of shape (28, 28)
Number of training examples : 10000 and each image is of shape (28, 28)

In [0]:

```
# if you observe the input shape its 2 dimensional vector
# for each image we have a (28*28) vector
# we will convert the (28*28) vector into single dimensional vector of 1 * 784

X_train = X_train.reshape(X_train.shape[0], X_train.shape[1]*X_train.shape[2])
X_test = X_test.reshape(X_test.shape[0], X_test.shape[1]*X_test.shape[2])
```

In [82]:

```
# after converting the input images from 3d to 2d vectors

print("Number of training examples :", X_train.shape[0], "and each image is of shape (%d)"%(X_train.shape[1]))
print("Number of training examples :", X_test.shape[0], "and each image is of shape (%d)"%(X_test.shape[1]))
```

Number of training examples : 60000 and each image is of shape (784)
Number of training examples : 10000 and each image is of shape (784)

In [83]:

```
# An example data point
```

[0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
	0	0	0	0	0	0	0	0	3	18	18	18	126	136	175	26	166	255	
247	127	0	0	0	0	0	0	0	0	0	0	0	0	0	30	36	94	154	
170	253	253	253	253	225	172	253	242	195	64	0	0	0	0	0	0	0	0	
	0	0	0	0	0	49	238	253	253	253	253	253	253	253	251	93	82		
82	56	39	0	0	0	0	0	0	0	0	0	0	0	0	0	18	219	253	
253	253	253	253	198	182	247	241	0	0	0	0	0	0	0	0	0	0	0	
	0	0	0	0	0	0	0	80	156	107	253	253	205	11	0	43	154		
	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
	0	14	1	154	253	90	0	0	0	0	0	0	0	0	0	0	0	0	
	0	0	0	0	0	0	0	0	0	0	0	0	0	139	253	190	2	0	
	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
	0	0	0	0	0	11	190	253	70	0	0	0	0	0	0	0	0	0	
	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	35	241	
225	160	108	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
	0	0	0	0	0	0	0	0	81	240	253	253	119	25	0	0	0	0	
	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
	0	0	45	186	253	253	150	27	0	0	0	0	0	0	0	0	0	0	
	0	0	0	0	0	0	0	0	0	0	0	0	0	16	93	252	253	187	
	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
	0	0	0	0	0	0	0	249	253	249	64	0	0	0	0	0	0	0	
	0	0	0	0	0	0	0	0	0	0	0	0	0	0	46	130	183	253	
253	207	2	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
	0	0	0	0	39	148	229	253	253	253	250	182	0	0	0	0	0	0	
	0	0	0	0	0														

```
# if we observe the above matrix each cell is having a value between 0-255
# before we move to apply machine learning algorithms lets try to normalize the data
#  $X \Rightarrow (X - X_{min}) / (X_{max} - X_{min}) = X / 255$ 

X_train = X_train/255
X_test = X_test/255
```

```
X_train = X_train/255
X_test = X_test/255
```

```
# example data point after normlizing
print(X_train[0])
```

[illegible]

[illegible]

[illegible]

In [86]:

```
# here we are having a class number for each image
print("Class label of first image :", y_train[0])

# lets convert this into a 10 dimensional vector
# ex: consider an image is 5 convert it into 5 => [0, 0, 0, 0, 0, 1, 0, 0, 0, 0]
# this conversion needed for MLPs

Y_train = np_utils.to_categorical(y_train, 10)
Y_test = np_utils.to_categorical(y_test, 10)

print("After converting the output into a vector : ",Y_train[0])
```

```
Class label of first image : 5
```

After converting the output into a vector : [0. 0. 0. 0. 0. 1. 0. 0. 0. 0.]

Softmax classifier

In [0]:

```
# https://keras.io/getting-started/sequential-model-guide/

# The Sequential model is a linear stack of layers.
# you can create a Sequential model by passing a list of layer instances to the constructor:

# model = Sequential([
#     Dense(32, input_shape=(784,)),
#     Activation('relu'),
#     Dense(10),
#     Activation('softmax')
# ])
```

```

# Activation( softmax ),
# ])

# You can also simply add layers via the .add() method:

# model = Sequential()
# model.add(Dense(32, input_dim=784))
# model.add(Activation('relu'))

###

# https://keras.io/layers/core/

# keras.layers.Dense(units, activation=None, use_bias=True, kernel_initializer='glorot_uniform',
# bias_initializer='zeros', kernel_regularizer=None, bias_regularizer=None,
# activity_regularizer=None,
# kernel_constraint=None, bias_constraint=None)

# Dense implements the operation: output = activation(dot(input, kernel) + bias) where
# activation is the element-wise activation function passed as the activation argument,
# kernel is a weights matrix created by the layer, and
# bias is a bias vector created by the layer (only applicable if use_bias is True).

# output = activation(dot(input, kernel) + bias) => y = activation(WT. X + b)

####

# https://keras.io/activations/

# Activations can either be used through an Activation layer, or through the activation argument s
upported by all forward layers:

# from keras.layers import Activation, Dense

# model.add(Dense(64))
# model.add(Activation('tanh'))

# This is equivalent to:
# model.add(Dense(64, activation='tanh'))

# there are many activation functions ar available ex: tanh, relu, softmax

from keras.models import Sequential
from keras.layers import Dense, Activation

```

In [0]:

```

# some model parameters

output_dim = 10
input_dim = X_train.shape[1]

batch_size = 128
nb_epoch = 20

```

MLP + Relu activation + ADAM with 2 hidden layers

In [89]:

```

# Multilayer perceptron

model_relu = Sequential()
model_relu.add(Dense(392, activation='relu', input_shape=(input_dim,)))
model_relu.add(Dense(196, activation='relu'))
model_relu.add(Dense(output_dim, activation='softmax'))

model_relu.summary()

```

Model: "sequential_20"

Layer (type)	Output Shape	Param #
dense_80 (Dense)	(None, 392)	397720

dense_00 (Dense)	(None, 32)	30720
dense_81 (Dense)	(None, 196)	77028
dense_82 (Dense)	(None, 10)	1970
=====		
Total params: 386,718		
Trainable params: 386,718		
Non-trainable params: 0		
=====		

In [90]:

```
model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

```
Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 6s 100us/step - loss: 0.2401 - acc: 0.9304 -
val_loss: 0.1109 - val_acc: 0.9654
Epoch 2/20
60000/60000 [=====] - 3s 53us/step - loss: 0.0881 - acc: 0.9730 -
val_loss: 0.0822 - val_acc: 0.9752
Epoch 3/20
60000/60000 [=====] - 3s 53us/step - loss: 0.0581 - acc: 0.9822 -
val_loss: 0.0708 - val_acc: 0.9787
Epoch 4/20
60000/60000 [=====] - 3s 53us/step - loss: 0.0388 - acc: 0.9879 -
val_loss: 0.0707 - val_acc: 0.9793
Epoch 5/20
60000/60000 [=====] - 3s 53us/step - loss: 0.0294 - acc: 0.9905 -
val_loss: 0.0715 - val_acc: 0.9777
Epoch 6/20
60000/60000 [=====] - 3s 53us/step - loss: 0.0215 - acc: 0.9932 -
val_loss: 0.0878 - val_acc: 0.9750
Epoch 7/20
60000/60000 [=====] - 3s 53us/step - loss: 0.0189 - acc: 0.9942 -
val_loss: 0.0772 - val_acc: 0.9795
Epoch 8/20
60000/60000 [=====] - 3s 53us/step - loss: 0.0164 - acc: 0.9947 -
val_loss: 0.0725 - val_acc: 0.9813
Epoch 9/20
60000/60000 [=====] - 3s 54us/step - loss: 0.0180 - acc: 0.9941 -
val_loss: 0.0691 - val_acc: 0.9817
Epoch 10/20
60000/60000 [=====] - 3s 53us/step - loss: 0.0115 - acc: 0.9961 -
val_loss: 0.0722 - val_acc: 0.9822
Epoch 11/20
60000/60000 [=====] - 3s 54us/step - loss: 0.0085 - acc: 0.9975 -
val_loss: 0.0861 - val_acc: 0.9810
Epoch 12/20
60000/60000 [=====] - 3s 53us/step - loss: 0.0137 - acc: 0.9957 -
val_loss: 0.0858 - val_acc: 0.9795
Epoch 13/20
60000/60000 [=====] - 3s 54us/step - loss: 0.0103 - acc: 0.9966 -
val_loss: 0.0717 - val_acc: 0.9833
Epoch 14/20
60000/60000 [=====] - 3s 53us/step - loss: 0.0091 - acc: 0.9971 -
val_loss: 0.0888 - val_acc: 0.9801
Epoch 15/20
60000/60000 [=====] - 3s 53us/step - loss: 0.0097 - acc: 0.9969 -
val_loss: 0.0929 - val_acc: 0.9794
Epoch 16/20
60000/60000 [=====] - 3s 53us/step - loss: 0.0106 - acc: 0.9965 -
val_loss: 0.0988 - val_acc: 0.9791
Epoch 17/20
60000/60000 [=====] - 3s 52us/step - loss: 0.0067 - acc: 0.9978 -
val_loss: 0.0868 - val_acc: 0.9824
Epoch 18/20
60000/60000 [=====] - 3s 53us/step - loss: 0.0060 - acc: 0.9980 -
val_loss: 0.0964 - val_acc: 0.9806
Epoch 19/20
60000/60000 [=====] - 3s 53us/step - loss: 0.0113 - acc: 0.9964 -
val_loss: 0.0958 - val_acc: 0.9809
```

```
Epoch 20/20
60000/60000 [=====] - 3s 53us/step - loss: 0.0074 - acc: 0.9978 -
val_loss: 0.0824 - val_acc: 0.9832
```

In [91]:

```
score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))


# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.08238437657678974

Test accuracy: 0.9832



In [92]:


```
w_after = model_relu.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w, color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')
```

```
plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



MLP + Relu activation + ADAM + Batch Normalization with 2 hidden layers

In [94]:

```
# Multilayer perceptron
from keras.layers.normalization import BatchNormalization
model_relu = Sequential()
model_relu.add(Dense(392, activation='relu', input_shape=(input_dim,)))
model_relu.add(BatchNormalization())
model_relu.add(Dense(196, activation='relu'))
model_relu.add(BatchNormalization())
model_relu.add(Dense(output_dim, activation='softmax'))

model_relu.summary()
```

Model: "sequential_22"

Layer (type)	Output Shape	Param #
dense_86 (Dense)	(None, 392)	307720
batch_normalization_32 (Batch Normalization)	(None, 392)	1568
dense_87 (Dense)	(None, 196)	77028
batch_normalization_33 (Batch Normalization)	(None, 196)	784
dense_88 (Dense)	(None, 10)	1970
Total params: 389,070		
Trainable params: 387,894		
Non-trainable params: 1,176		

In [95]:

```
model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 8s 139us/step - loss: 0.1823 - acc: 0.9451 - val_loss: 0.0967 - val_acc: 0.9684

Epoch 2/20

60000/60000 [=====] - 5s 83us/step - loss: 0.0722 - acc: 0.9776 - val_loss: 0.0856 - val_acc: 0.9706

Epoch 3/20


```

60000/60000 [=====] - 5s 88us/step - loss: 0.0473 - acc: 0.9853 -
val_loss: 0.0844 - val_acc: 0.9756
Epoch 4/20
60000/60000 [=====] - 5s 85us/step - loss: 0.0347 - acc: 0.9886 -
val_loss: 0.0809 - val_acc: 0.9756
Epoch 5/20
60000/60000 [=====] - 5s 83us/step - loss: 0.0266 - acc: 0.9918 -
val_loss: 0.0852 - val_acc: 0.9759
Epoch 6/20
60000/60000 [=====] - 5s 86us/step - loss: 0.0219 - acc: 0.9928 -
val_loss: 0.0974 - val_acc: 0.9721
Epoch 7/20
60000/60000 [=====] - 5s 85us/step - loss: 0.0219 - acc: 0.9925 -
val_loss: 0.0731 - val_acc: 0.9787
Epoch 8/20
60000/60000 [=====] - 5s 85us/step - loss: 0.0164 - acc: 0.9945 -
val_loss: 0.0741 - val_acc: 0.9793
Epoch 9/20
60000/60000 [=====] - 5s 83us/step - loss: 0.0162 - acc: 0.9945 -
val_loss: 0.0833 - val_acc: 0.9788
Epoch 10/20
60000/60000 [=====] - 5s 85us/step - loss: 0.0137 - acc: 0.9953 -
val_loss: 0.0745 - val_acc: 0.9776
Epoch 11/20
60000/60000 [=====] - 5s 84us/step - loss: 0.0136 - acc: 0.9954 -
val_loss: 0.0845 - val_acc: 0.9782
Epoch 12/20
60000/60000 [=====] - 5s 87us/step - loss: 0.0136 - acc: 0.9953 -
val_loss: 0.0834 - val_acc: 0.9797
Epoch 13/20
60000/60000 [=====] - 5s 86us/step - loss: 0.0117 - acc: 0.9958 -
val_loss: 0.0986 - val_acc: 0.9749
Epoch 14/20
60000/60000 [=====] - 5s 87us/step - loss: 0.0122 - acc: 0.9961 -
val_loss: 0.0802 - val_acc: 0.9813
Epoch 15/20
60000/60000 [=====] - 5s 88us/step - loss: 0.0085 - acc: 0.9972 -
val_loss: 0.0734 - val_acc: 0.9819
Epoch 16/20
60000/60000 [=====] - 5s 86us/step - loss: 0.0073 - acc: 0.9974 -
val_loss: 0.0825 - val_acc: 0.9819
Epoch 17/20
60000/60000 [=====] - 5s 90us/step - loss: 0.0103 - acc: 0.9968 -
val_loss: 0.0856 - val_acc: 0.9797
Epoch 18/20
60000/60000 [=====] - 5s 86us/step - loss: 0.0072 - acc: 0.9977 -
val_loss: 0.0763 - val_acc: 0.9825
Epoch 19/20
60000/60000 [=====] - 5s 86us/step - loss: 0.0073 - acc: 0.9975 -
val_loss: 0.0826 - val_acc: 0.9797
Epoch 20/20
60000/60000 [=====] - 5s 90us/step - loss: 0.0061 - acc: 0.9982 -
val_loss: 0.0850 - val_acc: 0.9826

```

In [96]:

```

score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, va
lida
tion_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

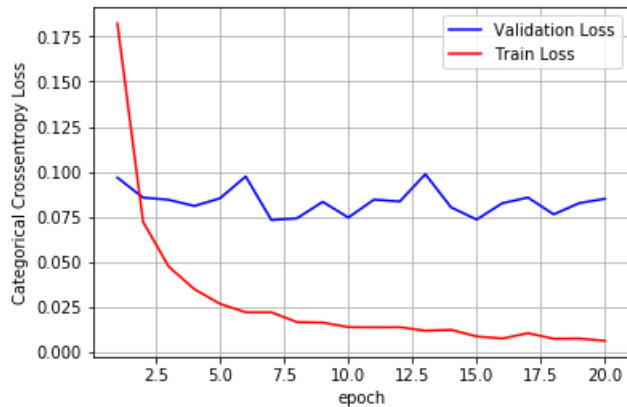
# loss : training loss

```

```
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.08495574688926236
Test accuracy: 0.9826



In [97]:


```
w_after = model_relu.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



MLP + Relu activation + ADAM + Dropout with 2 hidden layers

In [98]:

```
# Multilayer perceptron
from keras.layers import Dropout
model_relu = Sequential()
model_relu.add(Dense(392, activation='relu', input_shape=(input_dim,)))
model_relu.add(Dropout(0.5))
model_relu.add(Dense(196, activation='relu'))
model_relu.add(Dropout(0.5))
model_relu.add(Dense(output_dim, activation='softmax'))

model_relu.summary()
```

Model: "sequential_23"

Layer (type)	Output Shape	Param #
=====	=====	=====
dense_89 (Dense)	(None, 392)	307720
dropout_27 (Dropout)	(None, 392)	0
dense_90 (Dense)	(None, 196)	77028
dropout_28 (Dropout)	(None, 196)	0
dense_91 (Dense)	(None, 10)	1970
=====	=====	=====
Total params: 386,718		
Trainable params: 386,718		
Non-trainable params: 0		
=====		

In [99]:

```
model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

```
Epoch 1/20
60000/60000 [=====] - 7s 109us/step - loss: 0.4348 - acc: 0.8656 -
val_loss: 0.1421 - val_acc: 0.9565
Epoch 2/20
60000/60000 [=====] - 3s 58us/step - loss: 0.2000 - acc: 0.9408 -
val_loss: 0.1090 - val_acc: 0.9655
Epoch 3/20
60000/60000 [=====] - 3s 58us/step - loss: 0.1494 - acc: 0.9550 -
val_loss: 0.0946 - val_acc: 0.9688
Epoch 4/20
60000/60000 [=====] - 3s 57us/step - loss: 0.1313 - acc: 0.9604 -
val_loss: 0.0796 - val_acc: 0.9754
Epoch 5/20
60000/60000 [=====] - 4s 59us/step - loss: 0.1161 - acc: 0.9648 -
val_loss: 0.0798 - val_acc: 0.9763
Epoch 6/20
60000/60000 [=====] - 4s 59us/step - loss: 0.1044 - acc: 0.9688 -
val_loss: 0.0734 - val_acc: 0.9768
Epoch 7/20
60000/60000 [=====] - 3s 58us/step - loss: 0.0943 - acc: 0.9708 -
val_loss: 0.0693 - val_acc: 0.9788
Epoch 8/20
60000/60000 [=====] - 4s 59us/step - loss: 0.0892 - acc: 0.9726 -
val_loss: 0.0710 - val_acc: 0.9792
Epoch 9/20
60000/60000 [=====] - 4s 59us/step - loss: 0.0841 - acc: 0.9743 -
val_loss: 0.0687 - val_acc: 0.9799
Epoch 10/20
60000/60000 [=====] - 4s 59us/step - loss: 0.0752 - acc: 0.9765 -
val_loss: 0.0643 - val_acc: 0.9809
Epoch 11/20
60000/60000 [=====] - 4s 59us/step - loss: 0.0744 - acc: 0.9770 -
val_loss: 0.0600 - val_acc: 0.9811
Epoch 12/20
```

```

60000/60000 [=====] - 4s 59us/step - loss: 0.0699 - acc: 0.9785 -
val_loss: 0.0600 - val_acc: 0.9829
Epoch 13/20
60000/60000 [=====] - 4s 59us/step - loss: 0.0688 - acc: 0.9787 -
val_loss: 0.0625 - val_acc: 0.9820
Epoch 14/20
60000/60000 [=====] - 4s 58us/step - loss: 0.0645 - acc: 0.9797 -
val_loss: 0.0588 - val_acc: 0.9835
Epoch 15/20
60000/60000 [=====] - 4s 59us/step - loss: 0.0595 - acc: 0.9810 -
val_loss: 0.0643 - val_acc: 0.9827
Epoch 16/20
60000/60000 [=====] - 4s 59us/step - loss: 0.0589 - acc: 0.9814 -
val_loss: 0.0621 - val_acc: 0.9834
Epoch 17/20
60000/60000 [=====] - 4s 59us/step - loss: 0.0599 - acc: 0.9813 -
val_loss: 0.0652 - val_acc: 0.9823
Epoch 18/20
60000/60000 [=====] - 4s 60us/step - loss: 0.0564 - acc: 0.9817 -
val_loss: 0.0635 - val_acc: 0.9824
Epoch 19/20
60000/60000 [=====] - 4s 59us/step - loss: 0.0552 - acc: 0.9826 -
val_loss: 0.0623 - val_acc: 0.9828
Epoch 20/20
60000/60000 [=====] - 4s 59us/step - loss: 0.0536 - acc: 0.9831 -
val_loss: 0.0619 - val_acc: 0.9831

```

In [100]:

```

score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, va
lidaion_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy


# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.06188290925989168

Test accuracy: 0.9831



In [101]:


```
w_after = model_relu.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



MLP + Relu activation + ADAM + Batch Normalization + Dropout with 2 hidden layers

In [102]:

```
# Multilayer perceptron
model_relu = Sequential()
model_relu.add(Dense(392, activation='relu', input_shape=(input_dim,)))
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.5))
model_relu.add(Dense(196, activation='relu'))
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.5))
model_relu.add(Dense(output_dim, activation='softmax'))

model_relu.summary()
```

Model: "sequential_24"

Layer (type)	Output Shape	Param #
dense_92 (Dense)	(None, 392)	307720
batch_normalization_34 (Batch Normalization)	(None, 392)	1568
dropout_29 (Dropout)	(None, 392)	0
dense_93 (Dense)	(None, 196)	77028

batch_normalization_35 (Batch Normalization)	(None, 196)	784
dropout_30 (Dropout)	(None, 196)	0
dense_94 (Dense)	(None, 10)	1970
=====		
Total params: 389,070		
Trainable params: 387,894		
Non-trainable params: 1,176		

In [103]:

```
model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

```
Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 9s 152us/step - loss: 0.4204 - acc: 0.8732 -
val_loss: 0.1443 - val_acc: 0.9536
Epoch 2/20
60000/60000 [=====] - 5s 89us/step - loss: 0.2023 - acc: 0.9399 -
val_loss: 0.1098 - val_acc: 0.9676
Epoch 3/20
60000/60000 [=====] - 5s 90us/step - loss: 0.1625 - acc: 0.9502 -
val_loss: 0.0939 - val_acc: 0.9713
Epoch 4/20
60000/60000 [=====] - 5s 90us/step - loss: 0.1393 - acc: 0.9569 -
val_loss: 0.0881 - val_acc: 0.9740
Epoch 5/20
60000/60000 [=====] - 5s 90us/step - loss: 0.1219 - acc: 0.9622 -
val_loss: 0.0764 - val_acc: 0.9758
Epoch 6/20
60000/60000 [=====] - 5s 91us/step - loss: 0.1136 - acc: 0.9644 -
val_loss: 0.0725 - val_acc: 0.9768
Epoch 7/20
60000/60000 [=====] - 5s 90us/step - loss: 0.1054 - acc: 0.9673 -
val_loss: 0.0713 - val_acc: 0.9788
Epoch 8/20
60000/60000 [=====] - 5s 89us/step - loss: 0.0991 - acc: 0.9688 -
val_loss: 0.0674 - val_acc: 0.9794
Epoch 9/20
60000/60000 [=====] - 5s 90us/step - loss: 0.0930 - acc: 0.9703 -
val_loss: 0.0678 - val_acc: 0.9802
Epoch 10/20
60000/60000 [=====] - 5s 91us/step - loss: 0.0894 - acc: 0.9717 -
val_loss: 0.0672 - val_acc: 0.9802
Epoch 11/20
60000/60000 [=====] - 5s 91us/step - loss: 0.0825 - acc: 0.9745 -
val_loss: 0.0641 - val_acc: 0.9809
Epoch 12/20
60000/60000 [=====] - 5s 89us/step - loss: 0.0773 - acc: 0.9759 -
val_loss: 0.0595 - val_acc: 0.9811
Epoch 13/20
60000/60000 [=====] - 5s 91us/step - loss: 0.0779 - acc: 0.9750 -
val_loss: 0.0584 - val_acc: 0.9814
Epoch 14/20
60000/60000 [=====] - 6s 93us/step - loss: 0.0703 - acc: 0.9773 -
val_loss: 0.0598 - val_acc: 0.9822
Epoch 15/20
60000/60000 [=====] - 6s 92us/step - loss: 0.0690 - acc: 0.9781 -
val_loss: 0.0569 - val_acc: 0.9842
Epoch 16/20
60000/60000 [=====] - 5s 90us/step - loss: 0.0647 - acc: 0.9792 -
val_loss: 0.0593 - val_acc: 0.9813
Epoch 17/20
60000/60000 [=====] - 5s 90us/step - loss: 0.0646 - acc: 0.9799 -
val_loss: 0.0569 - val_acc: 0.9831
Epoch 18/20
60000/60000 [=====] - 5s 90us/step - loss: 0.0604 - acc: 0.9805 -
val_loss: 0.0553 - val_acc: 0.9841
Epoch 19/20
60000/60000 [=====] - 5s 91us/step - loss: 0.0595 - acc: 0.9807 -
```

```
val_loss: 0.0612 - val_acc: 0.9818
Epoch 20/20
60000/60000 [=====] - 6s 93us/step - loss: 0.0602 - acc: 0.9805 -
val_loss: 0.0569 - val_acc: 0.9829
```

In [104]:

```
score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))


# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

```
Test score: 0.05690854531452642
Test accuracy: 0.9829
```



In [105]:


```
w_after = model_relu.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w, color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')
```

```
plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



MLP + Relu activation + ADAM with 3 hidden layers

In [106]:

```
# Multilayer perceptron

model = Sequential()
model.add(Dense(516, activation='relu', input_shape=(input_dim,)))
model.add(Dense(258, activation='relu'))
model.add(Dense(50, activation='relu'))
model.add(Dense(output_dim, activation='softmax'))

model.summary()
```

Model: "sequential_25"

Layer (type)	Output Shape	Param #
dense_95 (Dense)	(None, 516)	405060
dense_96 (Dense)	(None, 258)	133386
dense_97 (Dense)	(None, 50)	12950
dense_98 (Dense)	(None, 10)	510
Total params: 551,906		
Trainable params: 551,906		
Non-trainable params: 0		

In [107]:

```
model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

```
Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 7s 120us/step - loss: 0.2391 - acc: 0.9292 -
val_loss: 0.1139 - val_acc: 0.9644
Epoch 2/20
60000/60000 [=====] - 4s 62us/step - loss: 0.0856 - acc: 0.9735 -
val_loss: 0.0925 - val_acc: 0.9718
Epoch 3/20
60000/60000 [=====] - 4s 63us/step - loss: 0.0548 - acc: 0.9827 -
val_loss: 0.0658 - val_acc: 0.9782
```



```

Epoch 4/20
60000/60000 [=====] - 4s 63us/step - loss: 0.0388 - acc: 0.9871 -
val_loss: 0.0800 - val_acc: 0.9747
Epoch 5/20
60000/60000 [=====] - 4s 63us/step - loss: 0.0295 - acc: 0.9904 -
val_loss: 0.0662 - val_acc: 0.9792
Epoch 6/20
60000/60000 [=====] - 4s 64us/step - loss: 0.0240 - acc: 0.9922 -
val_loss: 0.0681 - val_acc: 0.9798
Epoch 7/20
60000/60000 [=====] - 4s 64us/step - loss: 0.0203 - acc: 0.9935 -
val_loss: 0.0837 - val_acc: 0.9773
Epoch 8/20
60000/60000 [=====] - 4s 63us/step - loss: 0.0202 - acc: 0.9934 -
val_loss: 0.0769 - val_acc: 0.9787
Epoch 9/20
60000/60000 [=====] - 4s 63us/step - loss: 0.0183 - acc: 0.9940 -
val_loss: 0.0853 - val_acc: 0.9787
Epoch 10/20
60000/60000 [=====] - 4s 63us/step - loss: 0.0150 - acc: 0.9953 -
val_loss: 0.0779 - val_acc: 0.9806
Epoch 11/20
60000/60000 [=====] - 4s 62us/step - loss: 0.0127 - acc: 0.9957 -
val_loss: 0.0754 - val_acc: 0.9812
Epoch 12/20
60000/60000 [=====] - 4s 62us/step - loss: 0.0116 - acc: 0.9960 -
val_loss: 0.0871 - val_acc: 0.9800
Epoch 13/20
60000/60000 [=====] - 4s 63us/step - loss: 0.0097 - acc: 0.9968 -
val_loss: 0.0988 - val_acc: 0.9773
Epoch 14/20
60000/60000 [=====] - 4s 63us/step - loss: 0.0142 - acc: 0.9955 -
val_loss: 0.0917 - val_acc: 0.9796
Epoch 15/20
60000/60000 [=====] - 4s 63us/step - loss: 0.0102 - acc: 0.9964 -
val_loss: 0.0852 - val_acc: 0.9813
Epoch 16/20
60000/60000 [=====] - 4s 62us/step - loss: 0.0090 - acc: 0.9971 -
val_loss: 0.0885 - val_acc: 0.9819
Epoch 17/20
60000/60000 [=====] - 4s 63us/step - loss: 0.0065 - acc: 0.9978 -
val_loss: 0.0951 - val_acc: 0.9803
Epoch 18/20
60000/60000 [=====] - 4s 62us/step - loss: 0.0087 - acc: 0.9974 -
val_loss: 0.0978 - val_acc: 0.9802
Epoch 19/20
60000/60000 [=====] - 4s 63us/step - loss: 0.0088 - acc: 0.9974 -
val_loss: 0.0857 - val_acc: 0.9831
Epoch 20/20
60000/60000 [=====] - 4s 63us/step - loss: 0.0081 - acc: 0.9975 -
val_loss: 0.1060 - val_acc: 0.9781

```

In [108]:

```

score = model.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, va
lida
tion_data=(X_test, Y_test))


# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

```

```
vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.10599931993753521
Test accuracy: 0.9781



In [109]:

```
w_after = model.get_weights()

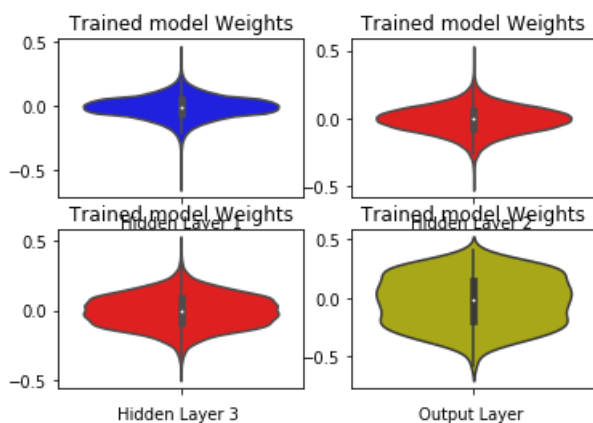
h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
out_w = w_after[6].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(2, 2, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(2, 2, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(2, 2, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='r')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(2, 2, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



MLP + Relu activation + ADAM + Batch Normalization with 3 hidden layers

In [110]:

```
# Multilayer perceptron
model = Sequential()
model.add(Dense(516, activation='relu', input_shape=(input_dim,)))
model.add(BatchNormalization())
model.add(Dense(258, activation='relu'))
model.add(BatchNormalization())
model.add(Dense(50, activation='relu'))
model.add(BatchNormalization())
model.add(Dense(output_dim, activation='softmax'))

model.summary()
```

Model: "sequential_26"

Layer (type)	Output Shape	Param #
dense_99 (Dense)	(None, 516)	405060
batch_normalization_36 (Batch Normalization)	(None, 516)	2064
dense_100 (Dense)	(None, 258)	133386
batch_normalization_37 (Batch Normalization)	(None, 258)	1032
dense_101 (Dense)	(None, 50)	12950
batch_normalization_38 (Batch Normalization)	(None, 50)	200
dense_102 (Dense)	(None, 10)	510
Total params: 555,202		
Trainable params: 553,554		
Non-trainable params: 1,648		

In [111]:

```
model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 11s 182us/step - loss: 0.1897 - acc: 0.9451 - val_loss: 0.0929 - val_acc: 0.9716

Epoch 2/20

60000/60000 [=====] - 7s 113us/step - loss: 0.0750 - acc: 0.9768 - val_loss: 0.1012 - val_acc: 0.9685

Epoch 3/20

60000/60000 [=====] - 7s 116us/step - loss: 0.0518 - acc: 0.9839 - val_loss: 0.1256 - val_acc: 0.9605

Epoch 4/20

60000/60000 [=====] - 7s 115us/step - loss: 0.0376 - acc: 0.9879 - val_loss: 0.0757 - val_acc: 0.9767

Epoch 5/20

60000/60000 [=====] - 7s 112us/step - loss: 0.0321 - acc: 0.9892 - val_loss: 0.0847 - val_acc: 0.9744

Epoch 6/20

60000/60000 [=====] - 7s 115us/step - loss: 0.0263 - acc: 0.9913 - val_loss: 0.0889 - val_acc: 0.9744

Epoch 7/20

60000/60000 [=====] - 7s 113us/step - loss: 0.0235 - acc: 0.9924 - val_loss: 0.0721 - val_acc: 0.9783

Epoch 8/20

60000/60000 [=====] - 7s 112us/step - loss: 0.0191 - acc: 0.9938 - val_loss: 0.0800 - val_acc: 0.9768

Epoch 9/20

```

60000/60000 [=====] - 7s 115us/step - loss: 0.0197 - acc: 0.9938 -
val_loss: 0.0853 - val_acc: 0.9757
Epoch 10/20
60000/60000 [=====] - 7s 115us/step - loss: 0.0159 - acc: 0.9951 -
val_loss: 0.0658 - val_acc: 0.9814
Epoch 11/20
60000/60000 [=====] - 7s 116us/step - loss: 0.0146 - acc: 0.9951 -
val_loss: 0.0695 - val_acc: 0.9819
Epoch 12/20
60000/60000 [=====] - 7s 114us/step - loss: 0.0126 - acc: 0.9955 -
val_loss: 0.0896 - val_acc: 0.9773
Epoch 13/20
60000/60000 [=====] - 7s 114us/step - loss: 0.0143 - acc: 0.9954 -
val_loss: 0.0747 - val_acc: 0.9802
Epoch 14/20
60000/60000 [=====] - 7s 119us/step - loss: 0.0147 - acc: 0.9949 -
val_loss: 0.0695 - val_acc: 0.9818
Epoch 15/20
60000/60000 [=====] - 7s 113us/step - loss: 0.0119 - acc: 0.9959 -
val_loss: 0.0661 - val_acc: 0.9816
Epoch 16/20
60000/60000 [=====] - 7s 114us/step - loss: 0.0099 - acc: 0.9965 -
val_loss: 0.0705 - val_acc: 0.9818
Epoch 17/20
60000/60000 [=====] - 7s 115us/step - loss: 0.0104 - acc: 0.9967 -
val_loss: 0.0873 - val_acc: 0.9785
Epoch 18/20
60000/60000 [=====] - 7s 114us/step - loss: 0.0110 - acc: 0.9962 -
val_loss: 0.0744 - val_acc: 0.9811
Epoch 19/20
60000/60000 [=====] - 7s 114us/step - loss: 0.0083 - acc: 0.9971 -
val_loss: 0.0649 - val_acc: 0.9832
Epoch 20/20
60000/60000 [=====] - 7s 114us/step - loss: 0.0065 - acc: 0.9978 -
val_loss: 0.0687 - val_acc: 0.9834

```

In [112]:

```

score = model.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, va
lidation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy


# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs


vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.06869542041149908

Test accuracy: 0.9834





In [113]:

```
w_after = model.get_weights()


h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
out_w = w_after[6].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(2, 2, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(2, 2, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(2, 2, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='r')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(2, 2, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



MLP + Relu activation + ADAM + Dropout with 3 hidden layers

In [114]:

```
# Multilayer perceptron
from keras.layers import Dropout
model = Sequential()
model.add(Dense(516, activation='relu', input_shape=(input_dim,)))
model.add(Dropout(0.5))
model.add(Dense(258, activation='relu'))
model.add(Dropout(0.5))
model.add(Dense(50, activation='relu'))
model.add(Dropout(0.5))
model.add(Dense(output_dim, activation='softmax'))
```

```
model.add(Dense(output_dim, activation='softmax'))
model.summary()
```

Model: "sequential_27"

Layer (type)	Output Shape	Param #
dense_103 (Dense)	(None, 516)	405060
dropout_31 (Dropout)	(None, 516)	0
dense_104 (Dense)	(None, 258)	133386
dropout_32 (Dropout)	(None, 258)	0
dense_105 (Dense)	(None, 50)	12950
dropout_33 (Dropout)	(None, 50)	0
dense_106 (Dense)	(None, 10)	510
Total params: 551,906		
Trainable params: 551,906		
Non-trainable params: 0		

In [115]:

```
model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

```
Epoch 1/20
60000/60000 [=====] - 8s 133us/step - loss: 0.6660 - acc: 0.7910 -
val_loss: 0.1822 - val_acc: 0.9470
Epoch 2/20
60000/60000 [=====] - 4s 69us/step - loss: 0.2760 - acc: 0.9281 -
val_loss: 0.1292 - val_acc: 0.9633
Epoch 3/20
60000/60000 [=====] - 4s 68us/step - loss: 0.2154 - acc: 0.9452 -
val_loss: 0.1217 - val_acc: 0.9668
Epoch 4/20
60000/60000 [=====] - 4s 68us/step - loss: 0.1818 - acc: 0.9538 -
val_loss: 0.1014 - val_acc: 0.9720
Epoch 5/20
60000/60000 [=====] - 4s 67us/step - loss: 0.1591 - acc: 0.9598 -
val_loss: 0.0941 - val_acc: 0.9744
Epoch 6/20
60000/60000 [=====] - 4s 69us/step - loss: 0.1517 - acc: 0.9611 -
val_loss: 0.0908 - val_acc: 0.9758
Epoch 7/20
60000/60000 [=====] - 4s 68us/step - loss: 0.1353 - acc: 0.9656 -
val_loss: 0.0834 - val_acc: 0.9777
Epoch 8/20
60000/60000 [=====] - 4s 68us/step - loss: 0.1271 - acc: 0.9673 -
val_loss: 0.0813 - val_acc: 0.9800
Epoch 9/20
60000/60000 [=====] - 4s 71us/step - loss: 0.1192 - acc: 0.9695 -
val_loss: 0.0813 - val_acc: 0.9792
Epoch 10/20
60000/60000 [=====] - 4s 71us/step - loss: 0.1106 - acc: 0.9707 -
val_loss: 0.0791 - val_acc: 0.9803
Epoch 11/20
60000/60000 [=====] - 4s 68us/step - loss: 0.1032 - acc: 0.9729 -
val_loss: 0.0788 - val_acc: 0.9807
Epoch 12/20
60000/60000 [=====] - 4s 68us/step - loss: 0.0978 - acc: 0.9745 -
val_loss: 0.0758 - val_acc: 0.9798
Epoch 13/20
60000/60000 [=====] - 4s 67us/step - loss: 0.0912 - acc: 0.9758 -
val_loss: 0.0775 - val_acc: 0.9825
Epoch 14/20
```

```

60000/60000 [=====] - 4s 68us/step - loss: 0.0951 - acc: 0.9749 -
val_loss: 0.0842 - val_acc: 0.9802
Epoch 15/20
60000/60000 [=====] - 4s 67us/step - loss: 0.0867 - acc: 0.9770 -
val_loss: 0.0786 - val_acc: 0.9814
Epoch 16/20
60000/60000 [=====] - 4s 68us/step - loss: 0.0849 - acc: 0.9773 -
val_loss: 0.0795 - val_acc: 0.9819
Epoch 17/20
60000/60000 [=====] - 4s 68us/step - loss: 0.0789 - acc: 0.9792 -
val_loss: 0.0798 - val_acc: 0.9828
Epoch 18/20
60000/60000 [=====] - 4s 68us/step - loss: 0.0780 - acc: 0.9789 -
val_loss: 0.0781 - val_acc: 0.9831
Epoch 19/20
60000/60000 [=====] - 4s 68us/step - loss: 0.0784 - acc: 0.9790 -
val_loss: 0.0779 - val_acc: 0.9830
Epoch 20/20
60000/60000 [=====] - 4s 68us/step - loss: 0.0751 - acc: 0.9799 -
val_loss: 0.0833 - val_acc: 0.9806

```

In [116]:

```

score = model.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, va
lvalidation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy


# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.08326539741690445

Test accuracy: 0.9806



In [117]:

```

w_after = model.get_weights()

```

```

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
out_w = w_after[6].flatten().reshape(-1,1)


fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(2, 2, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(2, 2, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(2, 2, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='r')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(2, 2, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```



MLP + Relu activation + ADAM + Batch Normalization + Dropout with 3 hidden layers

In [118]:

```

# Multilayer perceptron
model = Sequential()
model.add(Dense(516, activation='relu', input_shape=(input_dim,)))
model.add(BatchNormalization())
model.add(Dropout(0.5))
model.add(Dense(258, activation='relu'))
model.add(BatchNormalization())
model.add(Dropout(0.5))
model.add(Dense(50, activation='relu'))
model.add(BatchNormalization())
model.add(Dropout(0.5))
model.add(Dense(output_dim, activation='softmax'))

model.summary()

```

Model: "sequential_28"

Layer (type)	Output Shape	Param #
dense_107 (Dense)	(None, 516)	405060
batch_normalization_39 (Batch Normalization)	(None, 516)	2064
dropout_34 (Dropout)	(None, 516)	0

dense_108 (Dense)	(None, 258)	133386
batch_normalization_40 (Batch Normalization)	(None, 258)	1032
dropout_35 (Dropout)	(None, 258)	0
dense_109 (Dense)	(None, 50)	12950
batch_normalization_41 (Batch Normalization)	(None, 50)	200
dropout_36 (Dropout)	(None, 50)	0
dense_110 (Dense)	(None, 10)	510
=====		
Total params: 555,202		
Trainable params: 553,554		
Non-trainable params: 1,648		

In [119]:

```
model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

```
Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 12s 197us/step - loss: 0.6016 - acc: 0.8208 - val_loss: 0.1655 - val_acc: 0.9496
Epoch 2/20
60000/60000 [=====] - 7s 119us/step - loss: 0.2661 - acc: 0.9245 - val_loss: 0.1257 - val_acc: 0.9626
Epoch 3/20
60000/60000 [=====] - 7s 119us/step - loss: 0.2046 - acc: 0.9420 - val_loss: 0.0984 - val_acc: 0.9698
Epoch 4/20
60000/60000 [=====] - 7s 119us/step - loss: 0.1768 - acc: 0.9500 - val_loss: 0.0972 - val_acc: 0.9707
Epoch 5/20
60000/60000 [=====] - 7s 120us/step - loss: 0.1571 - acc: 0.9556 - val_loss: 0.0872 - val_acc: 0.9749
Epoch 6/20
60000/60000 [=====] - 7s 118us/step - loss: 0.1387 - acc: 0.9602 - val_loss: 0.0818 - val_acc: 0.9764
Epoch 7/20
60000/60000 [=====] - 7s 120us/step - loss: 0.1307 - acc: 0.9643 - val_loss: 0.0770 - val_acc: 0.9777
Epoch 8/20
60000/60000 [=====] - 7s 119us/step - loss: 0.1203 - acc: 0.9665 - val_loss: 0.0835 - val_acc: 0.9754
Epoch 9/20
60000/60000 [=====] - 7s 118us/step - loss: 0.1139 - acc: 0.9678 - val_loss: 0.0750 - val_acc: 0.9794
Epoch 10/20
60000/60000 [=====] - 7s 118us/step - loss: 0.1068 - acc: 0.9697 - val_loss: 0.0770 - val_acc: 0.9777
Epoch 11/20
60000/60000 [=====] - 7s 118us/step - loss: 0.1027 - acc: 0.9711 - val_loss: 0.0682 - val_acc: 0.9813
Epoch 12/20
60000/60000 [=====] - 7s 122us/step - loss: 0.0972 - acc: 0.9721 - val_loss: 0.0657 - val_acc: 0.9820
Epoch 13/20
60000/60000 [=====] - 7s 125us/step - loss: 0.0932 - acc: 0.9741 - val_loss: 0.0663 - val_acc: 0.9816
Epoch 14/20
60000/60000 [=====] - 7s 118us/step - loss: 0.0881 - acc: 0.9749 - val_loss: 0.0702 - val_acc: 0.9800
Epoch 15/20
60000/60000 [=====] - 7s 121us/step - loss: 0.0854 - acc: 0.9759 - val_loss: 0.0700 - val_acc: 0.9813
Epoch 16/20
60000/60000 [=====] - 7s 120us/step - loss: 0.0825 - acc: 0.9761 - val_loss: 0.0666 - val_acc: 0.9816
```

```

val_loss: 0.0666 - val_acc: 0.9819
Epoch 17/20
60000/60000 [=====] - 7s 118us/step - loss: 0.0795 - acc: 0.9774 -
val_loss: 0.0618 - val_acc: 0.9824
Epoch 18/20
60000/60000 [=====] - 7s 121us/step - loss: 0.0732 - acc: 0.9786 -
val_loss: 0.0560 - val_acc: 0.9850
Epoch 19/20
60000/60000 [=====] - 7s 121us/step - loss: 0.0737 - acc: 0.9780 -
val_loss: 0.0628 - val_acc: 0.9818
Epoch 20/20
60000/60000 [=====] - 7s 120us/step - loss: 0.0730 - acc: 0.9781 -
val_loss: 0.0592 - val_acc: 0.9829

```

In [120]:

```

score = model.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, va
lvalidation_data=(X_test, Y_test))


# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.059158047910808816
Test accuracy: 0.9829



In [121]:

```

w_after = model.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
out_w = w_after[6].flatten().reshape(-1,1)

fig = plt.figure()
plt.plot(h1_w, label='h1 weights', color='red')
plt.plot(h2_w, label='h2 weights', color='blue')
plt.plot(h3_w, label='h3 weights', color='green')
plt.plot(out_w, label='out weights', color='purple')
plt.legend()
plt.show()


```

```
plt.title("weight matrices arter model trained")
plt.subplot(2, 2, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(2, 2, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(2, 2, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='r')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(2, 2, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



MLP + Relu activation + ADAM with 5 hidden layers

In [122]:

```
# Multilayer perceptron

model2 = Sequential()
model2.add(Dense(645, activation='relu', input_shape=(input_dim,)))
model2.add(Dense(510, activation='relu'))
model2.add(Dense(387, activation='relu'))
model2.add(Dense(252, activation='relu'))
model2.add(Dense(129, activation='relu'))
model2.add(Dense(output_dim, activation='softmax'))

model2.summary()
```

Model: "sequential_29"

Layer (type)	Output Shape	Param #
=====	=====	=====
dense_111 (Dense)	(None, 645)	506325
dense_112 (Dense)	(None, 510)	329460
dense_113 (Dense)	(None, 387)	197757
dense_114 (Dense)	(None, 252)	97776
dense_115 (Dense)	(None, 129)	32637
dense_116 (Dense)	(None, 10)	1300
=====	=====	=====
Total params: 1,165,255		
Trainable params: 1,165,255		
Non-trainable params: 0		

In [123]:

```
model2.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model2.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

```
Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 9s 156us/step - loss: 0.2283 - acc: 0.9299 -
val_loss: 0.1040 - val_acc: 0.9682
Epoch 2/20
60000/60000 [=====] - 5s 84us/step - loss: 0.0901 - acc: 0.9723 -
val_loss: 0.0889 - val_acc: 0.9734
Epoch 3/20
60000/60000 [=====] - 5s 83us/step - loss: 0.0650 - acc: 0.9804 -
val_loss: 0.0925 - val_acc: 0.9730
Epoch 4/20
60000/60000 [=====] - 5s 83us/step - loss: 0.0486 - acc: 0.9850 -
val_loss: 0.0845 - val_acc: 0.9755
Epoch 5/20
60000/60000 [=====] - 5s 84us/step - loss: 0.0412 - acc: 0.9877 -
val_loss: 0.0768 - val_acc: 0.9798
Epoch 6/20
60000/60000 [=====] - 5s 84us/step - loss: 0.0338 - acc: 0.9894 -
val_loss: 0.0764 - val_acc: 0.9788
Epoch 7/20
60000/60000 [=====] - 5s 84us/step - loss: 0.0281 - acc: 0.9915 -
val_loss: 0.0843 - val_acc: 0.9802
Epoch 8/20
60000/60000 [=====] - 5s 84us/step - loss: 0.0286 - acc: 0.9916 -
val_loss: 0.0764 - val_acc: 0.9800
Epoch 9/20
60000/60000 [=====] - 5s 82us/step - loss: 0.0226 - acc: 0.9937 -
val_loss: 0.0831 - val_acc: 0.9791
Epoch 10/20
60000/60000 [=====] - 5s 83us/step - loss: 0.0252 - acc: 0.9927 -
val_loss: 0.0793 - val_acc: 0.9817
Epoch 11/20
60000/60000 [=====] - 5s 83us/step - loss: 0.0170 - acc: 0.9948 -
val_loss: 0.0869 - val_acc: 0.9811
Epoch 12/20
60000/60000 [=====] - 5s 83us/step - loss: 0.0199 - acc: 0.9941 -
val_loss: 0.0877 - val_acc: 0.9815
Epoch 13/20
60000/60000 [=====] - 5s 89us/step - loss: 0.0150 - acc: 0.9956 -
val_loss: 0.0908 - val_acc: 0.9818
Epoch 14/20
60000/60000 [=====] - 5s 89us/step - loss: 0.0141 - acc: 0.9958 -
val_loss: 0.0829 - val_acc: 0.9832
Epoch 15/20
60000/60000 [=====] - 5s 83us/step - loss: 0.0162 - acc: 0.9952 -
val_loss: 0.1045 - val_acc: 0.9793
Epoch 16/20
60000/60000 [=====] - 5s 84us/step - loss: 0.0181 - acc: 0.9954 -
val_loss: 0.0815 - val_acc: 0.9821
Epoch 17/20
60000/60000 [=====] - 5s 84us/step - loss: 0.0111 - acc: 0.9969 -
val_loss: 0.0987 - val_acc: 0.9810
Epoch 18/20
60000/60000 [=====] - 5s 84us/step - loss: 0.0109 - acc: 0.9971 -
val_loss: 0.0920 - val_acc: 0.9825
Epoch 19/20
60000/60000 [=====] - 5s 83us/step - loss: 0.0155 - acc: 0.9957 -
val_loss: 0.0920 - val_acc: 0.9808
Epoch 20/20
60000/60000 [=====] - 5s 84us/step - loss: 0.0160 - acc: 0.9958 -
val_loss: 0.0916 - val_acc: 0.9831
```

In [124]:

```
score = model2.evaluate(X_test, Y_test, verbose=0)
```

```

print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, va
lidation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy


# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.09155124295547798

Test accuracy: 0.9831



In [125]:

```

w_after = model2.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
h4_w = w_after[6].flatten().reshape(-1,1)
h5_w = w_after[8].flatten().reshape(-1,1)
out_w = w_after[10].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(2, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(2, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(2, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='r')
plt.xlabel('Hidden Layer 3 ')

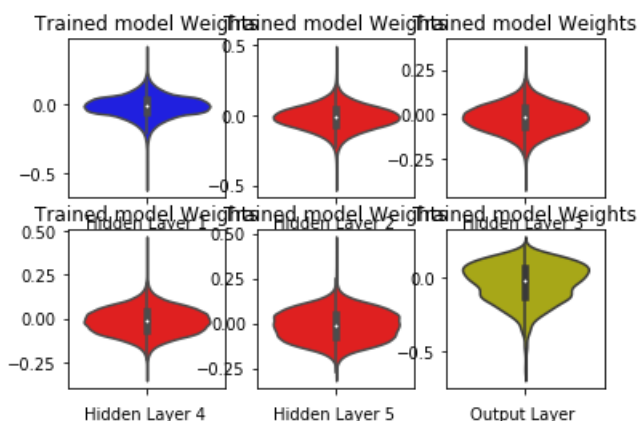
plt.subplot(2, 3, 4)

```

```
plt.subplot(2, 3, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h4_w, color='r')
plt.xlabel('Hidden Layer 4 ')

plt.subplot(2, 3, 5)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h5_w, color='r')
plt.xlabel('Hidden Layer 5 ')

plt.subplot(2, 3, 6)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



MLP + Relu activation + ADAM + Batch Normalization with 5 hidden layers

In [126]:

```
# Multilayer perceptron

model2 = Sequential()
model2.add(Dense(645, activation='relu', input_shape=(input_dim,)))
model2.add(BatchNormalization())
model2.add(Dense(510, activation='relu'))
model2.add(BatchNormalization())
model2.add(Dense(387, activation='relu'))
model2.add(BatchNormalization())
model2.add(Dense(252, activation='relu'))
model2.add(BatchNormalization())
model2.add(Dense(129, activation='relu'))
model2.add(BatchNormalization())
model2.add(Dense(output_dim, activation='softmax'))

model2.summary()
```

Model: "sequential_30"

Layer (type)	Output Shape	Param #
=====	=====	=====
dense_117 (Dense)	(None, 645)	506325
dense_118 (Dense)	(None, 510)	329460
dense_119 (Dense)	(None, 387)	197757
dense_120 (Dense)	(None, 252)	97776
dense_121 (Dense)	(None, 129)	32637
dense_122 (Dense)	(None, 10)	1300
=====	=====	=====
Total params:	1,165,255	
Trainable params:	1,165,255	
Non-trainable params:	0	

In [127]:

```
model2.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model2.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

```
Epoch 1/20
60000/60000 [=====] - 10s 161us/step - loss: 0.2261 - acc: 0.9313 - val_loss: 0.1190 - val_acc: 0.9636
Epoch 2/20
60000/60000 [=====] - 5s 84us/step - loss: 0.0934 - acc: 0.9722 - val_loss: 0.0854 - val_acc: 0.9738
Epoch 3/20
60000/60000 [=====] - 5s 84us/step - loss: 0.0620 - acc: 0.9813 - val_loss: 0.0900 - val_acc: 0.9750
Epoch 4/20
60000/60000 [=====] - 5s 84us/step - loss: 0.0478 - acc: 0.9849 - val_loss: 0.0846 - val_acc: 0.9777
Epoch 5/20
60000/60000 [=====] - 5s 84us/step - loss: 0.0415 - acc: 0.9870 - val_loss: 0.0788 - val_acc: 0.9783
Epoch 6/20
60000/60000 [=====] - 5s 84us/step - loss: 0.0323 - acc: 0.9900 - val_loss: 0.0745 - val_acc: 0.9801
Epoch 7/20
60000/60000 [=====] - 5s 84us/step - loss: 0.0297 - acc: 0.9909 - val_loss: 0.0865 - val_acc: 0.9784
Epoch 8/20
60000/60000 [=====] - 5s 85us/step - loss: 0.0273 - acc: 0.9919 - val_loss: 0.0814 - val_acc: 0.9788
Epoch 9/20
60000/60000 [=====] - 5s 84us/step - loss: 0.0203 - acc: 0.9941 - val_loss: 0.0993 - val_acc: 0.9797
Epoch 10/20
60000/60000 [=====] - 5s 84us/step - loss: 0.0191 - acc: 0.9942 - val_loss: 0.0722 - val_acc: 0.9824
Epoch 11/20
60000/60000 [=====] - 5s 84us/step - loss: 0.0225 - acc: 0.9935 - val_loss: 0.0863 - val_acc: 0.9791
Epoch 12/20
60000/60000 [=====] - 5s 85us/step - loss: 0.0147 - acc: 0.9954 - val_loss: 0.1001 - val_acc: 0.9784
Epoch 13/20
60000/60000 [=====] - 5s 84us/step - loss: 0.0176 - acc: 0.9952 - val_loss: 0.0930 - val_acc: 0.9802
Epoch 14/20
60000/60000 [=====] - 5s 84us/step - loss: 0.0152 - acc: 0.9958 - val_loss: 0.0885 - val_acc: 0.9812
Epoch 15/20
60000/60000 [=====] - 5s 84us/step - loss: 0.0134 - acc: 0.9960 - val_loss: 0.0859 - val_acc: 0.9820
Epoch 16/20
60000/60000 [=====] - 5s 84us/step - loss: 0.0143 - acc: 0.9958 - val_loss: 0.0807 - val_acc: 0.9826
Epoch 17/20
60000/60000 [=====] - 5s 85us/step - loss: 0.0151 - acc: 0.9957 - val_loss: 0.0772 - val_acc: 0.9835
Epoch 18/20
60000/60000 [=====] - 5s 83us/step - loss: 0.0125 - acc: 0.9966 - val_loss: 0.0913 - val_acc: 0.9832
Epoch 19/20
60000/60000 [=====] - 5s 84us/step - loss: 0.0115 - acc: 0.9968 - val_loss: 0.1136 - val_acc: 0.9807
Epoch 20/20
60000/60000 [=====] - 5s 84us/step - loss: 0.0126 - acc: 0.9966 - val_loss: 0.0911 - val_acc: 0.9830
```

In [128]:

```
score = model2.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
```

```

print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, va
lvalidation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy


# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.09110382709081515

Test accuracy: 0.983



In [129]:

```

w_after = model2.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
h4_w = w_after[6].flatten().reshape(-1,1)
h5_w = w_after[8].flatten().reshape(-1,1)
out_w = w_after[10].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(2, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(2, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(2, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='r')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(2, 3, 4)
plt.title("Trained model Weights")

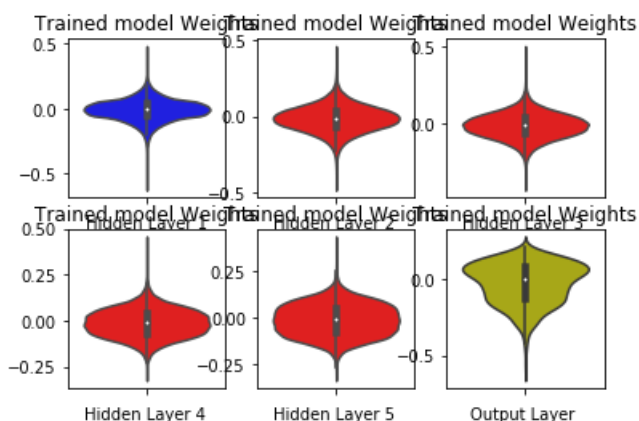
```



```
plt.title('Trained model Weights')
ax = sns.violinplot(y=h4_w, color='r')
plt.xlabel('Hidden Layer 4')

plt.subplot(2, 3, 5)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h5_w, color='r')
plt.xlabel('Hidden Layer 5')

plt.subplot(2, 3, 6)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w, color='y')
plt.xlabel('Output Layer')
plt.show()
```



MLP + Relu activation + ADAM + Dropout with 5 hidden layers

In [130]:

```
# Multilayer perceptron

model2 = Sequential()
model2.add(Dense(645, activation='relu', input_shape=(input_dim,)))
model2.add(Dropout(0.5))
model2.add(Dense(510, activation='relu'))
model2.add(Dropout(0.5))
model2.add(Dense(387, activation='relu'))
model2.add(Dropout(0.5))
model2.add(Dense(252, activation='relu'))
model2.add(Dropout(0.5))
model2.add(Dense(129, activation='relu'))
model2.add(Dropout(0.5))
model2.add(Dense(output_dim, activation='softmax'))

model2.summary()
```

Model: "sequential_31"

Layer (type)	Output Shape	Param #
dense_123 (Dense)	(None, 645)	506325
dropout_37 (Dropout)	(None, 645)	0
dense_124 (Dense)	(None, 510)	329460
dropout_38 (Dropout)	(None, 510)	0
dense_125 (Dense)	(None, 387)	197757
dropout_39 (Dropout)	(None, 387)	0
dense_126 (Dense)	(None, 252)	97776
dropout_40 (Dropout)	(None, 252)	0
dense_127 (Dense)	(None, 129)	32637

dropout_41 (Dropout)	(None, 129)	0
dense_128 (Dense)	(None, 10)	1300
=====		
Total params: 1,165,255		
Trainable params: 1,165,255		
Non-trainable params: 0		
=====		

In [131]:

```
model2.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model2.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

```
Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 10s 169us/step - loss: 0.7789 - acc: 0.7449 - val_loss: 0.2111 - val_acc: 0.9400
Epoch 2/20
60000/60000 [=====] - 5s 89us/step - loss: 0.2748 - acc: 0.9308 - val_loss: 0.1316 - val_acc: 0.9643
Epoch 3/20
60000/60000 [=====] - 5s 91us/step - loss: 0.2149 - acc: 0.9460 - val_loss: 0.1159 - val_acc: 0.9700
Epoch 4/20
60000/60000 [=====] - 5s 90us/step - loss: 0.1863 - acc: 0.9542 - val_loss: 0.1058 - val_acc: 0.9734
Epoch 5/20
60000/60000 [=====] - 5s 90us/step - loss: 0.1616 - acc: 0.9590 - val_loss: 0.0959 - val_acc: 0.9740
Epoch 6/20
60000/60000 [=====] - 5s 89us/step - loss: 0.1554 - acc: 0.9623 - val_loss: 0.0964 - val_acc: 0.9754
Epoch 7/20
60000/60000 [=====] - 5s 89us/step - loss: 0.1433 - acc: 0.9642 - val_loss: 0.0905 - val_acc: 0.9761
Epoch 8/20
60000/60000 [=====] - 5s 89us/step - loss: 0.1350 - acc: 0.9661 - val_loss: 0.0860 - val_acc: 0.9781
Epoch 9/20
60000/60000 [=====] - 5s 90us/step - loss: 0.1239 - acc: 0.9689 - val_loss: 0.0839 - val_acc: 0.9795
Epoch 10/20
60000/60000 [=====] - 5s 89us/step - loss: 0.1222 - acc: 0.9701 - val_loss: 0.0880 - val_acc: 0.9788
Epoch 11/20
60000/60000 [=====] - 5s 89us/step - loss: 0.1113 - acc: 0.9724 - val_loss: 0.0834 - val_acc: 0.9790
Epoch 12/20
60000/60000 [=====] - 5s 91us/step - loss: 0.1117 - acc: 0.9722 - val_loss: 0.0823 - val_acc: 0.9804
Epoch 13/20
60000/60000 [=====] - 5s 90us/step - loss: 0.1068 - acc: 0.9738 - val_loss: 0.0854 - val_acc: 0.9814
Epoch 14/20
60000/60000 [=====] - 5s 90us/step - loss: 0.0999 - acc: 0.9755 - val_loss: 0.0770 - val_acc: 0.9820
Epoch 15/20
60000/60000 [=====] - 5s 90us/step - loss: 0.1020 - acc: 0.9755 - val_loss: 0.0801 - val_acc: 0.9818
Epoch 16/20
60000/60000 [=====] - 5s 90us/step - loss: 0.0956 - acc: 0.9765 - val_loss: 0.0867 - val_acc: 0.9800
Epoch 17/20
60000/60000 [=====] - 5s 90us/step - loss: 0.0921 - acc: 0.9771 - val_loss: 0.0854 - val_acc: 0.9803
Epoch 18/20
60000/60000 [=====] - 5s 90us/step - loss: 0.0927 - acc: 0.9778 - val_loss: 0.0835 - val_acc: 0.9818
Epoch 19/20
60000/60000 [=====] - 5s 89us/step - loss: 0.0887 - acc: 0.9784 - val_loss: 0.0726 - val_acc: 0.9827
Epoch 20/20
```

```
Epoch 20/20
60000/60000 [=====] - 5s 90us/step - loss: 0.0886 - acc: 0.9788 -
val_loss: 0.0790 - val_acc: 0.9830
```

In [132]:

```
score = model2.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, va
lvalidation_data=(X_test, Y_test))


# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.0790262889877944

Test accuracy: 0.983



In [133]:

```
w_after = model2.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
h4_w = w_after[6].flatten().reshape(-1,1)
h5_w = w_after[8].flatten().reshape(-1,1)
out_w = w_after[10].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(2, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w, color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(2, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
```

```

ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')


plt.subplot(2, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='r')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(2, 3, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h4_w, color='r')
plt.xlabel('Hidden Layer 4 ')

plt.subplot(2, 3, 5)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h5_w, color='r')
plt.xlabel('Hidden Layer 5 ')

plt.subplot(2, 3, 6)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```



MLP + Relu activation + ADAM + Batch Normalization + Dropout with 5 hidden layers

In [134]:

```

# Multilayer perceptron
model2 = Sequential()
model2.add(Dense(645, activation='relu', input_shape=(input_dim,)))
model2.add(BatchNormalization())
model2.add(Dropout(0.5))
model2.add(Dense(510, activation='relu'))
model2.add(BatchNormalization())
model2.add(Dropout(0.5))
model2.add(Dense(387, activation='relu'))
model2.add(BatchNormalization())
model2.add(Dropout(0.5))
model2.add(Dense(252, activation='relu'))
model2.add(BatchNormalization())
model2.add(Dropout(0.5))
model2.add(Dense(129, activation='relu'))
model2.add(BatchNormalization())
model2.add(Dropout(0.5))
model2.add(Dense(output_dim, activation='softmax'))

model2.summary()

```

Model: "sequential_32"

Layer (type)	Output Shape	Param #
dense_129 (Dense)	(None, 645)	506325
batch_normalization_47 (Batch Normalization)	(None, 645)	2580

dropout_42 (Dropout)	(None, 645)	0
dense_130 (Dense)	(None, 510)	329460
batch_normalization_48 (Batch Normalization)	(None, 510)	2040
dropout_43 (Dropout)	(None, 510)	0
dense_131 (Dense)	(None, 387)	197757
batch_normalization_49 (Batch Normalization)	(None, 387)	1548
dropout_44 (Dropout)	(None, 387)	0
dense_132 (Dense)	(None, 252)	97776
batch_normalization_50 (Batch Normalization)	(None, 252)	1008
dropout_45 (Dropout)	(None, 252)	0
dense_133 (Dense)	(None, 129)	32637
batch_normalization_51 (Batch Normalization)	(None, 129)	516
dropout_46 (Dropout)	(None, 129)	0
dense_134 (Dense)	(None, 10)	1300
=====		
Total params: 1,172,947		
Trainable params: 1,169,101		
Non-trainable params: 3,846		

In [135]:

```
model2.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model2.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

```
Epoch 1/20
60000/60000 [=====] - 17s 282us/step - loss: 0.7552 - acc: 0.7723 - val_loss: 0.2059 - val_acc: 0.9419
Epoch 2/20
60000/60000 [=====] - 11s 179us/step - loss: 0.2782 - acc: 0.9222 - val_loss: 0.1457 - val_acc: 0.9567
Epoch 3/20
60000/60000 [=====] - 11s 181us/step - loss: 0.2141 - acc: 0.9400 - val_loss: 0.1085 - val_acc: 0.9701
Epoch 4/20
60000/60000 [=====] - 11s 179us/step - loss: 0.1807 - acc: 0.9496 - val_loss: 0.1031 - val_acc: 0.9707
Epoch 5/20
60000/60000 [=====] - 11s 181us/step - loss: 0.1630 - acc: 0.9545 - val_loss: 0.0924 - val_acc: 0.9741
Epoch 6/20
60000/60000 [=====] - 11s 184us/step - loss: 0.1512 - acc: 0.9581 - val_loss: 0.0910 - val_acc: 0.9740
Epoch 7/20
60000/60000 [=====] - 11s 184us/step - loss: 0.1378 - acc: 0.9611 - val_loss: 0.0820 - val_acc: 0.9771
Epoch 8/20
60000/60000 [=====] - 11s 181us/step - loss: 0.1276 - acc: 0.9634 - val_loss: 0.0829 - val_acc: 0.9763
Epoch 9/20
60000/60000 [=====] - 11s 180us/step - loss: 0.1252 - acc: 0.9647 - val_loss: 0.0776 - val_acc: 0.9778
Epoch 10/20
60000/60000 [=====] - 11s 181us/step - loss: 0.1178 - acc: 0.9669 - val_loss: 0.0713 - val_acc: 0.9806
Epoch 11/20
60000/60000 [=====] - 11s 180us/step - loss: 0.1093 - acc: 0.9693 - val_loss: 0.0723 - val_acc: 0.9792
Epoch 12/20
60000/60000 [=====] - 11s 180us/step - loss: 0.1050 - acc: 0.9704 - val_loss: 0.0704 - val_acc: 0.9792
```

```

60000/60000 [-----] - 11s 160us/step - loss: 0.1000 - acc: 0.9704 - val_l
oss: 0.0728 - val_acc: 0.9810
Epoch 13/20
60000/60000 [=====] - 11s 179us/step - loss: 0.1001 - acc: 0.9720 - val_l
oss: 0.0667 - val_acc: 0.9813
Epoch 14/20
60000/60000 [=====] - 11s 179us/step - loss: 0.0985 - acc: 0.9716 - val_l
oss: 0.0736 - val_acc: 0.9795
Epoch 15/20
60000/60000 [=====] - 11s 180us/step - loss: 0.0896 - acc: 0.9747 - val_l
oss: 0.0674 - val_acc: 0.9817
Epoch 16/20
60000/60000 [=====] - 11s 180us/step - loss: 0.0913 - acc: 0.9740 - val_l
oss: 0.0634 - val_acc: 0.9821
Epoch 17/20
60000/60000 [=====] - 11s 180us/step - loss: 0.0865 - acc: 0.9764 - val_l
oss: 0.0641 - val_acc: 0.9821
Epoch 18/20
60000/60000 [=====] - 11s 182us/step - loss: 0.0851 - acc: 0.9755 - val_l
oss: 0.0628 - val_acc: 0.9832
Epoch 19/20
60000/60000 [=====] - 11s 182us/step - loss: 0.0815 - acc: 0.9766 - val_l
oss: 0.0652 - val_acc: 0.9816
Epoch 20/20
60000/60000 [=====] - 11s 180us/step - loss: 0.0776 - acc: 0.9781 - val_l
oss: 0.0647 - val_acc: 0.9832

```

In [136]:

```

score = model2.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, va
lidation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy


# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.06474938120327424

Test accuracy: 0.9832



In [137]:

```
w_after = model2.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
h4_w = w_after[6].flatten().reshape(-1,1)
h5_w = w_after[8].flatten().reshape(-1,1)
out_w = w_after[10].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(2, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')


plt.subplot(2, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(2, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='r')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(2, 3, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h4_w, color='r')
plt.xlabel('Hidden Layer 4 ')

plt.subplot(2, 3, 5)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h5_w, color='r')
plt.xlabel('Hidden Layer 5 ')

plt.subplot(2, 3, 6)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



Observations

All the below architectures are implemented with fixing **relu** as **activation function** and **adam** as **optimization technique**. We found that whatever might be the architecture, the model was overfitting when implemented with neither dropout nor batch normalization and with just batch normalization. The model is converging fine when implemented either with just dropout or with both dropout and batch normalization.

In [3]:

```
from prettytable import PrettyTable
```

```

from prettytable import PrettyTable

x = PrettyTable()

x.field_names = ["Architecture", "NoOfNeurons/Layer", "BatchNormalization", "Dropouts", "Accuracy"]

x.add_row(["2 Hidden Layers", "784-392-196-10", "No", "No", "98.32%"])
x.add_row(["2 Hidden Layers", "784-392-196-10", "Yes", "No", "98.26%"])
x.add_row(["2 Hidden Layers", "784-392-196-10", "No", "Yes", "98.31%"])
x.add_row(["2 Hidden Layers", "784-392-196-10", "Yes", "Yes", "98.29%"])
x.add_row(["3 Hidden Layers", "784-516-258-50-10", "No", "No", "97.81%"])
x.add_row(["3 Hidden Layers", "784-516-258-50-10", "Yes", "No", "98.34%"])
x.add_row(["3 Hidden Layers", "784-516-258-50-10", "No", "Yes", "98.06%"])
x.add_row(["3 Hidden Layers", "784-516-258-50-10", "Yes", "Yes", "98.29%"])
x.add_row(["5 Hidden Layers", "784-645-510-387-252-129-10", "No", "No", "98.31%"])
x.add_row(["5 Hidden Layers", "784-645-510-387-252-129-10", "Yes", "No", "98.30%"])
x.add_row(["5 Hidden Layers", "784-645-510-387-252-129-10", "No", "Yes", "98.30%"])
x.add_row(["5 Hidden Layers", "784-645-510-387-252-129-10", "Yes", "Yes", "98.32%"])

print(x)

```

Architecture	NoOfNeurons/Layer	BatchNormalization	Dropouts	Accuracy
2 Hidden Layers	784-392-196-10	No	No	98.32%
2 Hidden Layers	784-392-196-10	Yes	No	98.26%
2 Hidden Layers	784-392-196-10	No	Yes	98.31%
2 Hidden Layers	784-392-196-10	Yes	Yes	98.29%
3 Hidden Layers	784-516-258-50-10	No	No	97.81%
3 Hidden Layers	784-516-258-50-10	Yes	No	98.34%
3 Hidden Layers	784-516-258-50-10	No	Yes	98.06%
3 Hidden Layers	784-516-258-50-10	Yes	Yes	98.29%
5 Hidden Layers	784-645-510-387-252-129-10	No	No	98.31%
5 Hidden Layers	784-645-510-387-252-129-10	Yes	No	98.30%
5 Hidden Layers	784-645-510-387-252-129-10	No	Yes	98.30%
5 Hidden Layers	784-645-510-387-252-129-10	Yes	Yes	98.32%